

A decorative graphic on the left side of the slide, consisting of a network of white lines and small circles on a dark blue background, resembling a circuit board or a neural network.

KUBERNETES TASKS

DEVOPS MASTER ARCHITECT WORKSHOP

TASK 1:HOW TO CREATE KUBERNETES CLUSTER – MINIKUBE WAY

- Activity: Install Kubernetes cluster in minikube
- Definition of Done:
 - Minikube installed
 - Create a GitHub repository for Kubernetes task. Create a subfolder and upload the relevant screenshot.(Terminal screenshot of version of the Minikube Kubernetes installed)

#	Steps	Commands
1	Clone the Kubernetes repository	git clone https://github.com/AnjuMeleth/DevOpsMasterKubernetesTasks.git
2	Run the installation script in 1_installation sub folder	sh minikube_install.sh
3	Get the kubectl version	sudo kubectl version --client=true
4	Get the minikube version	sudo minikube version

TASK 2:HOW TO START A MINIKUBE CLUSTER

- Activity: Start a minikube cluster on Ubuntu 18.04 system
- Definition of Done:
 - Create a subfolder and upload the relevant screenshot.(Minikube started)

#	Steps	Commands
1	Start a Minikube cluster	<code>sudo minikube start</code>
2	Throws an error "Unable to start VM as this computer doesn't have VT-X/AMD-v enabled"	
3	Start Minikube with no driver	<code>sudo minikube start --vm-driver=none</code>
4	Throws an error ""docker": executable file not found in \$PATH". Install docker	<code>sudo apt-get install docker.io</code>
5	Start the minikube once again	<code>sudo minikube start --vm-driver=none</code>

TASK 3:HOW TO DEPLOY AN APPLICATION ON MINIKUBE CLUSTER USING PODS

#	Steps	Commands
1	Run a pod on the minikube cluster	<code>sudo kubectl run nginx --image=nginx --port=80</code>
2	View the pods	<code>sudo kubectl get pods</code>

- Activity: Deploy a Nginx application on to Minikube cluster as a pod
- Definition of Done:
 - Pod created .
 - Create a subfolder and upload the relevant screenshot.(List of pods running)

TASK 4:HOW TO ACCESS AN APPLICATION RUNNING IN A POD FROM OUTSIDE

- Activity: Create a NodePort service that helps to connect to pod created in the last task.
- Definition of Done:
 - Web server accessible at a port
 - Create a subfolder and upload the relevant screenshot.(Nginx service accessible at the NodePort)

#	Steps	Commands
1	Create a NodePort service	<code>sudo kubectl expose pod nginx -type=NodePort</code>
2	List the services	<code>sudo kubectl get services</code>
3	Access the server from the browser	<code>http://<Ip address of the system>:port</code>
4	Delete the service	<code>sudo kubectl delete svc nginx</code>
5	Stop the Minikube cluster	<code>sudo minikube stop</code>
6	Delete the Minikube cluster	<code>sudo minikube delete</code>
7	Delete the config file	<code>sudo rm -rf ~/.kube/config</code>
8	Delete the kubernetes	<code>rm -rf /etc/kubernetes</code>

TASK 5:HOW TO ESTABLISH PASSWORDLESS CONNECTIVITY ACROSS INSTANCES GOING TO BE USED AS MANAGER AND WORKER NODES FOR COMMUNICATION

- Activity: Establish passwordless connectivity by using ssh agent
- Definition of Done:
 - Able to connect to worker nodes from the manager without password
 - No screenshot to be uploaded

#	Steps	Commands
1	Transfer the .pem file to Kubernetes manager node using Filezilla	
2	Stop all firewall	sudo service ufw stop
3	Start ssh	sudo service ssh start
4	Add the ssh details	eval `ssh-agent -s` sudo chmod 400 JenkinsMumbai.pem ssh-add JenkinsMumbai.pem
5	Access the node system from master	ssh ubuntu@<IP address>
6	Stop all firewall	sudo service ufw stop

TASK 6 :HOW TO BRING UP A KUBERNETES MASTER IN A PRODUCTION CLUSTER USING KUBEADM

- Activity: Create a Kubernetes master using Kubeadm installation
- Definition of Done:
 - Kubernetes master up and running.
 - Create a subfolder and upload the relevant screenshot.(No screenshots needed)

#	Steps	Commands
1	Modify the hostname in Kubernetes manager node	<code>sudo vi /etc/hostname</code>
2	Enter the Kubernetes master and node details in /etc/hosts file of the master	<code>sudo vi /etc/hosts</code>
3	Run the installation script	<code>sh kubeadm_install.sh</code>
4	Edit the file /etc/systemd/system/kubelet.service.d/10-kubeadm.conf with the content	<code>sudo nano /etc/systemd/system/kubelet.service.d/10-kubeadm.conf</code> Environment="cgroup-driver=systemd/cgroup-driver=cgroupfs"
5	Initialise the Kubernetes cluster by advertising the master IP address. Make a note of the output details	<code>sudo kubeadm init --pod-network-cidr=10.244.0.0/16</code>

Task 6 :how to bring up a Kubernetes master in a production cluster using kubeadm Contd...

6	Run the noted commands	<pre>mkdir -p \$HOME/.kube sudo cp -i /etc/kubernetes/admin.conf \$HOME/.kube/config sudo chown \$(id -u):\$(id - \$HOME/.kube/config</pre>
7	List the pods	<pre>kubectl get pods -o wide --all- namespaces</pre>
8	Deploy a pod network	<pre>kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/2140ac876ef134e0ed5af15c65e414cf26827915/Documentation/kube-flannel.yml</pre>
9	List the pods	<pre>kubectl get pods -o wide --all- namespaces</pre>

TASK 7:HOW TO ADD A WORKER NODE TO KUBERNETES CLUSTER

- Activity: Add a worker node Ubuntu 18.04 to Kubernetes cluster
- Definition of Done:
 - Node added to cluster
 - Create a subfolder and upload the relevant screenshot.(List of nodes added to the cluster)

#	Steps	Commands
1	Make sure you have the right hostname details in /etc/hosts file and /etc/hostname files	<pre>sudo vi /etc/hostname sudo vi /etc/hosts</pre>
2	Run the installation script	<pre>sh kubeadm_install.sh</pre>
3	Edit the file /etc/systemd/system/kubelet.service.d/10-kubeadm.conf with the content	<pre>sudo nano /etc/systemd/system/kubelet.se vice.d/10-kubeadm.conf</pre> <pre>Environment="cgroup- driver=systemd/cgroup- driver=cgroupfs"</pre>
4	Run the join token command received during the kubeadm init command	<pre>sudo kubeadm join <>--token <token></pre>
5	Get the list of nodes in the Kubernetes master	<pre>kubectl get nodes</pre>

TASK 8: HOW TO DEPLOY AN APPLICATION POD ON KUBERNETES CLUSTER AND ACCESS IT

- Activity: Apply a pod of nginx application to the Kubernetes cluster and access it by a NodePort service
- Definition of Done:
 - Nginx up , running and accessible at a NodePort
 - Create a subfolder and upload the relevant screenshot.(Nginx web page on the worker node)

#	Steps	Commands
1	Run a pod on the cluster	<code>sudo kubectl run nginx --image=nginx --port=80</code>
2	View the pods	<code>sudo kubectl get pods</code>
3	Create a NodePort service	<code>sudo kubectl expose pod nginx --type=NodePort</code>
4	List the services	<code>sudo kubectl get svc</code>
5	Access the server from the browser	<code>http://< Correct Ip address of the system>:port</code>
6	Get more details about the pods	<code>kubectl describe pods nginx</code>
7	Delete the service	<code>sudo kubectl delete svc nginx</code>
8	Delete pods	<code>kubectl delete pods nginx</code>



TASK 9 : HOW TO DEPLOY A POD ON KUBERNETES CLUSTER AND ACCESS IT USING KUBECTL

- Activity: Deploy apache application on the Kubernetes cluster
- Definition of Done:
 - Apache up , running and accessible at NodePort.
 - Create a subfolder and upload the relevant screenshot.(Apache web page on the worker node)

#	Steps	Commands
1	Run a pod on the cluster	<code>sudo kubectl run apache --image=httpd --port=80</code>
2	View the pods	<code>sudo kubectl get pods</code>
3	Create a NodePort service	<code>sudo kubectl expose pod apache --type=NodePort</code>
4	List the services	<code>sudo kubectl get svc</code>
5	Access the server from the browser	<code>http://< Correct Ip address of the system>:port</code>
6	Get more details about the pods	<code>kubectl describe pods apache</code>

TASK 10 : HOW TO DEPLOY A POD ON KUBERNETES CLUSTER USING YML

- Activity: Deploy nginx on the Kubernetes cluster using yml
- Definition of Done:
 - Nginx up , running and accessible at NodePort.
 - Create a subfolder and upload the relevant screenshot.(No screenshot to be uploaded)

#	Steps	Commands
1	Create nginx_pod.yml	vi nginx_pod.yml
2	Apply the yml	kubectl apply -f nginx_pod.yml
3	Get the list of pods running	kubectl get pods
4	Delete the pods	kubectl delete pods mypod

TASK 11:HOW TO CREATE A REPLICATION CONTROLLER

- Activity: Create a replication controller for Nginx application so that it maintains the number of instances of the same pod running
- Definition of Done:
 - Replication controller created.
 - Create a subfolder and upload the relevant screenshot.(Terminal screenshot of list of replication controllers created)

#	Steps	Commands
1	Edit the nginx_rc.yml file in 11_rc folder	vi nginx_rc.yml
2	Apply the yml file	kubectl apply -f nginx_rc.yml
3	List the replication controller	kubectl get rc
4	Get the list of pods	kubectl get pods
5	Get the details about the replication controller created	kubectl describe replicationcontrollers/nginx
6	To list only the pods under the replication controller	pods=\$(kubectl get pods --selector=app=nginx --output=jsonpath={.items..metadata.name}) echo \$pods

TASK 12:HOW TO SCALE AN APPLICATION IN A KUBERNETES CLUSTER

- Activity: Scale the Nginx application from the previous example to run 5 instances to cater to a greater demand

#	Steps	Commands
1	Scale the application to run 5 instances	kubectl scale replicationcontrollers/nginx --replicas=5
2	List the rc	kubectl get rc
3	List the pods	kubectl get pods

- Definition of Done:
 - Create a subfolder and upload the relevant screenshot.(List of 5 pods running)

TASK 13 : HOW TO ACCESS AN APPLICATION FROM OUTSIDE

- Activity: Create a service for the Nginx application running in the replication controller in previous tasks

- Definition of Done:

- Create a subfolder and upload the relevant screenshot.()

#	Steps	Commands
1	Edit the nginx_service.yml in 13_service folder	vi nginx_service.yml
2	Apply the yml	kubectl apply -f nginx_service.yml
3	Get the list of services	kubectl get svc
4	Access the Nginx application in port 8080	http://<Ip address>:port

TASK 14 : HOW DOES SELF HEALING WORK IN REPLICATION CONTROLLER

- Activity: Delete all pods in the replication controller in previous example

- Definition of Done:
 - New pods gets created. No screenshots to be uploaded.

#	Steps	Commands
1	Get the list of rc	kubectrl get rc
2	Get the list of services	kubectrl get svc
3	Get the list of the pods running	kubectrl get pods
4	Delete all pods	kubectrl delete po --all
5	Get the list of the pods running	kubectrl get pods

TASK 15 : HOW DOES THE CLIENT PODS GET TO KNOW THE IP ADDRESS AND PORT OF SERVICES(SERVICE DISCOVERY)

- Activity: All the services in the Kubernetes cluster has an Ip assigned and this could be seen in the pods created . If the service is created before the pods recreate the pods once again.
- Definition of Done:
 - Environment variables could be seen in the pod.
 - Create a subfolder and upload the relevant screenshot.(Terminal screenshot of environment variables in a pod)

#	Steps	Commands
1	Get the list of rc	kubectrl get rc
2	Get the list of services	kubectrl get svc
3	Get the list of the pods running	kubectrl get pods
4	View the environment variables in any one pod	kubectrl exec <pod name> env

TASK 16: HOW TO CREATE A DEPLOYMENT IN KUBERNETES CLUSTER AND ACCESS IT FROM OUTSIDE

- Activity: Create a deployment that has 5 instances of nginx application running. Then access the nginx service from a NodePort
- Definition of Done:
 - Deployment and service created.
 - Create a subfolder and upload the relevant screenshot.(Terminal screenshot of the list of deployments in the Kubernetes cluster)

#	Steps	Commands
1	Edit the nginx_deployent.yml file in 16_deployment folder	vi nginx_deployment.yml
2	Apply the yml file	kubectl apply -f nginx_deployment.yml
3	List the replication controller	kubectl get deployment
4	Get the list of pods	kubectl get pods
5	Get more details about the deployment	kubectl describe deployment/nginx-deployment
6	To list only the pods under the replication controller	pods=\$(kubectl get pods --selector=app=nginx1 --output=jsonpath={.items..metadata.name}) echo \$pods
7	Create a service	vi nginx_service.yml

TASK 16 : CONTD...

8	Apply the service yml file	kubectl apply -f nginx_service.yml
9	List the services	kubectl get svc
10	Access the Nginx application	http://<IpAddress>:port

TASK 17 : HOW TO DO ROLLING UPDATE IN KUBERNETES

- **Activity:** Create a deployment that has 3 instance of a Node js(version1) application running. We need to update it with a newer version version 2
- **Definition of Done:**
 - Pods run with newer version
 - Create a subfolder and upload the relevant screenshot.(Terminal screenshot showing the version of the image used in pods of replication controller)

#	Steps	Commands
1	Edit the file kubia.yml in 17_rolling update folder	vi kubia.yml
2	Apply	kubectl apply -f kubia.yml
3	Get the pods	kubectl get pods -l app=kubia
4	Modify the image of the pods	kubectl set image deployment kubia nodejs=luksa/kubia:v2
5	Get the list of pods	kubectl get pods -l app=kubia
6	Get the details of the deployment	kubectl describe deployment kubia

TASK 18 :HOW TO DO A ROLL BACK IN KUBERNETES

- Activity: Update the deployment created in the previous task with an error image. Then perform a roll back
- Definition of Done:
 - All the pods uses the previous version of the image.
 - Create a subfolder and upload the relevant screenshot.(Terminal screenshot showing the rollback status)

#	Steps	Commands
1	Update the image used in the deployment kubia to a new version	kubectrl set image deployment kubia nodejs=luksa/kubia:v3
2	Access the server	http://<Ipaddress>:port
3	Refresh the webserver for few times so that we get error message	
4	Roll back to previous version	kubectrl rollout undo deployment kubia

TASK 19 : HOW TO DO CANARY DEPLOYMENTS IN KUBERNETES

#	Steps	Commands
1	Update the image of the deployment with Version 4	kubectl set image deployment kubia nodejs=luksa/kubia:v4
2	Pause the deployment	kubectl rollout pause deployment kubia
3	Resume the deployment	kubectl rollout resume deployment kubia

- Activity: Update the deployment created with a new version. Pause the update . Resume it
- Definition of Done:
 - Canary deployment is performed
 - Create a subfolder and upload the relevant screenshot.(Terminal screenshot showing the new pod with new image)

#	Steps	Commands
1	Goto the folder 19_canary and apply the yml configuration to the cluster	kubectl apply -f newkubia.yml
2	Get the services from the nodeport address	http://<ip address>:Nodeport

TASK 20 :HOW TO INSTALL KUBERNETES DASHBOARD

- Activity: Install Kubernetes GUI
- Definition of Done:
 - Kubernetes dashboard is visible
 - Create a subfolder and upload the relevant screenshot.(screenshot of Kubernetes dashboard)

#	Steps	Commands
1	Apply dashboard yml available	<code>kubectl apply -f https://raw.githubusercontent.com/kubernetes/dashboard/v2.0.0-beta4/aio/deploy/recommended.yaml</code>
2	Edit the services and put the type as NodePort	<code>kubectl edit services kubernetes-dashboard -n kubernetes-dashboard</code>
3	Run <code>kubectl proxy</code> and open another terminal	<code>kubectl proxy &</code>
4	Get the services list and note down the NodePort	<code>kubectl get svc --all-namespaces</code>
5	Open any browser other than chrome. Go to Advanced and accept risk and continue	<code>https://<ipaddress>:port</code>
6	Create a service account	<code>kubectl create serviceaccount dashboard -n default</code>
7	Create cluster role binding	<code>kubectl create clusterrolebinding dashboard-admin -n default --clusterrole=cluster-admin --serviceaccount=default:dashboard</code>

TASK 20:CONTD...

#	Steps	Commands
8	Get the secret token	<code>kubectl get secret \$(kubectl get serviceaccount dashboard -o jsonpath="{.secrets[0].name}") -o jsonpath="{.data.token}" base64 --decode</code>
9	Copy and paste the token in the dashboard	

TASK 21 :HOW TO DEPLOY SPRING PET CLINIC APPLICATION

- Activity: Deploy the spring pet clinic application on to the Kubernetes cluster
- Definition of Done:
 - Spring pet clinic application is accessible at the NodePort
 - Create a subfolder and upload the relevant screenshot.(screenshot of spring pet clinic application)

#	Steps	Commands
1	Apply the yml in 21_springpetclinic folder	kubectl apply -f petclinic.yml
2	Get the Nodeport port from the service	kubectl get svc
2	Access the application	http://<ipaddress>:port